

DOCUMENT CONTROL

Hiring Scorecard — Sarah Chen

Structured evaluation sheet for the Senior ML Engineer hiring workflow.

Document ID	HR-003
Owner	Talent Acquisition
Candidate	Sarah Chen
Role	Senior ML Engineer
Review Date	2024-03-15

Synthetic corporate document for demo, evaluation, and retrieval workflows only. All parties, people, systems, repositories, and events are fictional.

Score Summary

Dimension	Weight	Rating	Weighted Score	Evidence
LLM / transformer expertise	20	5 / 5	20	Hands-on transformer work, model optimization, and conference publications.
Production ML systems at scale	20	5 / 5	20	Designed high-throughput inference systems and MLOps automation.
Technical leadership	15	4 / 5	12	Led a five-person ML pod and mentored junior engineers.
RAG / retrieval depth	15	3 / 5	9	Strong adjacent skills; direct retrieval-system evidence limited in CV.
Compliance / governance awareness	10	2 / 5	4	No direct SOC2 / HIPAA ownership shown.
Edge / on-prem deployment	10	2 / 5	4	Mostly cloud-native deployments.
Cross-functional collaboration	10	4 / 5	8	Repeated evidence of product and research collaboration.

Overall recommendation: 87 / 100 — Strong Hire. Advance to final technical interview with emphasis on RAG architecture and regulated-workload instincts.

Decision Signals

- Meets all three must-have areas already evidenced in the CV: LLM expertise, large-scale production ML, and technical/team leadership.
- Primary gaps are specific to retrieval architecture and regulated workload exposure, both considered trainable if fundamentals are strong.
- No blocking risk was found in communication, ownership, or seniority signals.