

Sarah Chen

Senior ML Engineer • San Francisco, CA • sarah.chen@candidate-mail.example • +1 (415) 555-0172

SUMMARY

Senior ML engineer with 8 years of experience across research, model optimization, and production systems. Strong record in transformer-based workloads, large-scale inference, and technical leadership in fast-moving product environments.

CORE STRENGTHS

- Transformer architectures, instruction tuning, evaluation design, and training-cost optimisation.
- Production ML systems with 10M+ daily predictions and modern MLOps tooling.
- Technical leadership for small ML teams, mentorship, and cross-functional delivery.
- Research communication, publication writing, and experiment-driven iteration.

EXPERIENCE

Aurora AI Systems — Senior ML Engineer • 2022 — Present

Led RLHF and evaluation pipeline improvements for enterprise language-assistance products; reduced training and evaluation costs by 34%; partnered with product and infra teams to roll out production model updates with audit-friendly change logs.

Atlas Data Robotics — ML Engineer • 2020 — 2022

Built data quality scoring and annotation automation for large-scale ML workflows; supported high-throughput inference systems serving more than 10M daily predictions; collaborated closely with product and operations stakeholders.

Northstar Research Labs — Research Engineer • 2018 — 2020

Contributed to transformer-model experiments, published two peer-reviewed papers on attention mechanisms, and mentored interns on evaluation pipelines and reproducibility practices.

EDUCATION

M.S. Computer Science, Pacific Coast University • 2018

B.S. Electrical Engineering and Computer Science, Bay City Institute • 2016

PUBLICATIONS AND LEADERSHIP

- Five peer-reviewed conference papers across model evaluation and efficient transformer training.
- Led a five-person ML pod during a product-scale rollout at Aurora AI Systems.
- Mentored three interns and two junior engineers on experimentation, tooling, and model release process.

TOOLS AND SYSTEMS

PyTorch, JAX, Kubeflow, MLflow, Weights & Biases, Kubernetes, feature stores, experiment tracking, CI for ML, observability-driven rollout.