

DOCUMENT CONTROL

Senior ML Engineer Role Brief

Hiring brief for Meridian’s AI Platform Engineering team.

Document ID	HR-002
Owner	AI Platform Engineering
Hiring Manager	Nora Patel
Open Date	2024-02-26
Classification	Internal Hiring

Synthetic corporate document for demo, evaluation, and retrieval workflows only. All parties, people, systems, repositories, and events are fictional.

Role Purpose

Meridian is hiring a Senior ML Engineer to lead production-grade LLM and retrieval systems used in knowledge search, analyst assistance, and enterprise automation. The role partners with Platform, Security, Product, and Infrastructure teams.

Core Responsibilities

- Design and operate production LLM and retrieval pipelines with measurable latency, relevance, and reliability goals.
- Lead architecture decisions for fine-tuning, evaluation, observability, and model-serving patterns.
- Mentor engineers and influence cross-functional technical roadmaps.
- Partner with Security and Compliance on AI-specific controls, auditability, and regulated-workload requirements.

Role-Fit Criteria

Requirement	Type	Notes
LLM / transformer expertise	Must-have	Model architecture, fine-tuning, and evaluation depth expected.
Production ML systems at scale	Must-have	Experience with large-scale inference and MLOps operations.
Team leadership experience	Must-have	Technical leadership and mentoring track record required.
RAG / retrieval system design	Must-have	Vector retrieval, chunking, ranking, and relevance trade-offs.
Regulatory compliance (SOC2, HIPAA)	Nice-to-have	Direct experience is valuable but trainable.
Edge / on-prem deployment	Nice-to-have	Useful for constrained environments and customer-specific architectures.
Cross-functional collaboration	Nice-to-have	Ability to work across Product, Infra, Security, and Analytics teams.